

Optimizations in Processing

Here are **10 key optimizations** you made to enhance processing efficiency:

1. Delta Lake for Storage & Query Optimization

- Used Delta Lake for storing Go vulnerability incidents, ensuring ACID transactions and efficient time travel.
- Delta caching improves read performance significantly over raw Parquet.

2. Partitioning on Year, Month, and Day

- Partitioning by year, month, and day allows pruning of unnecessary data, reducing query scan time for time-based queries.

3. Use of Databricks Autoloader

- Used **Databricks Autoloader** (`cloudFiles`) to efficiently ingest files from **ADLS**, reducing overhead compared to manual batch ingestion.

4. Efficient Data Transformation with Spark SQL & PySpark

- Applied **distributed transformations** using **Spark SQL** instead of traditional row-wise operations, improving efficiency.
- Used **broadcast joins** for small tables to optimize join performance.

5. Delta Lake Merge for SCD Type 2 Implementation

- Implemented **MERGE INTO** for SCD Type 2 in `dim_packages`, preventing unnecessary overwrites and reducing data duplication.

6. Compression & File Format Optimization

- Used **Parquet format with Snappy compression**, reducing storage size and improving read speeds compared to CSV or JSON.

7. Z-Ordering for Faster Queries

- Used **Z-ORDER BY (modified_dt, published_dt)** to co-locate data for faster query execution when filtering on timestamps.

8. **Adaptive Query Execution (AQE) in Databricks**

- Leveraged **AQE** to dynamically optimize joins and shuffle partitions at runtime, reducing unnecessary memory consumption.

9. **Auto-Scaling Compute for Cost Efficiency**

- Used **Databricks Auto-Scaling clusters** with **spot instances**, optimizing costs while handling varying workloads.

10. **Incremental Processing with Change Data Capture (CDC)**

- Instead of full reprocessing, only **modified and new records** are processed using **timestamp-based CDC**, reducing unnecessary computations.